

# JSC 270 - LECTURE 9

## EMBEDDING

<https://jsc270.github.io/>



# DIMENSIONALITY REDUCTION

How would you define it?



# DIMENSIONALITY REDUCTION

What applications can you think of?



# DIMENSIONALITY REDUCTION

Which techniques do you know?



## DIMENSIONALITY REDUCTION VS EMBEDDING

High-dimensional data often lies on or near a much lower dimensional manifold.

We reduce dimension or *embed* data into the manifold to get the *embedding*, in other words a low dimensional representation

# USEFUL CONCEPTS: MANIFOLD

**Manifold** – a collection of points forming a certain kind of set, such as those of a topologically closed surface in three or more dimensions (such topological space that locally resembles Euclidean space near each point)

Swiss roll



sphere



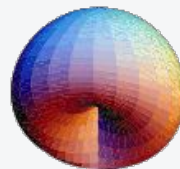
torus



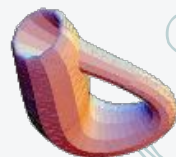
double torus



cross surface



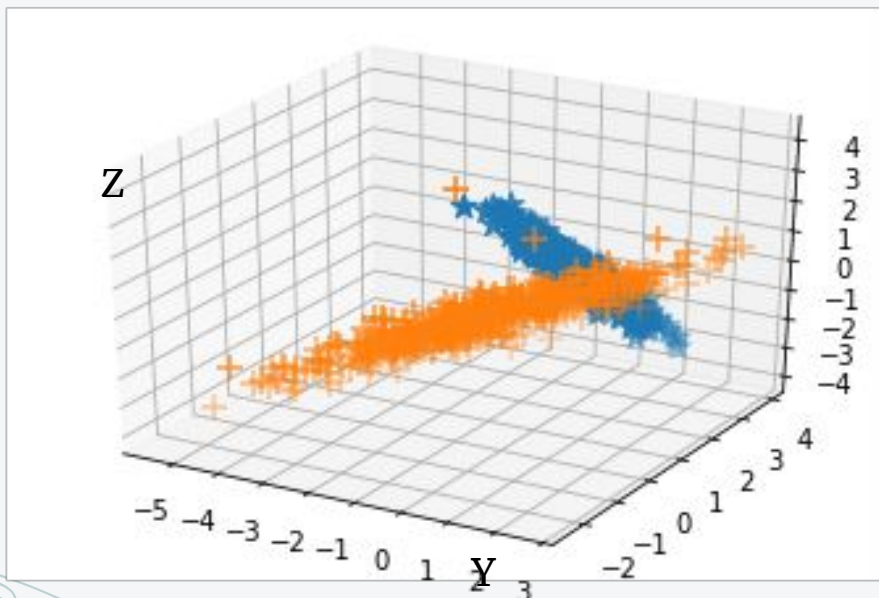
Klein bottle



# DATA

Number of samples  $N = 1000$

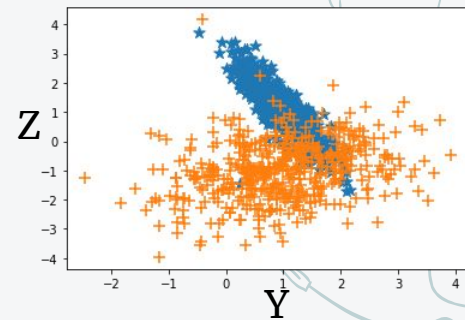
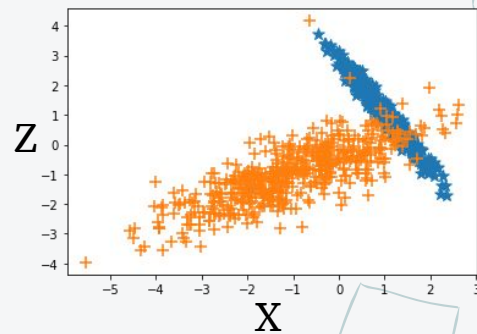
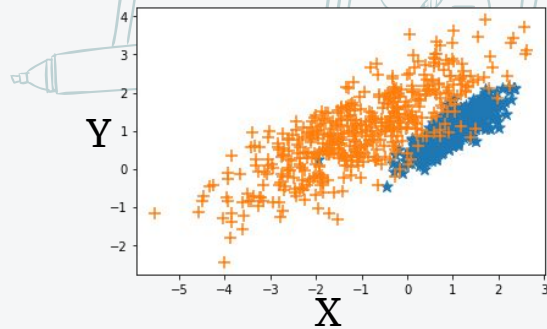
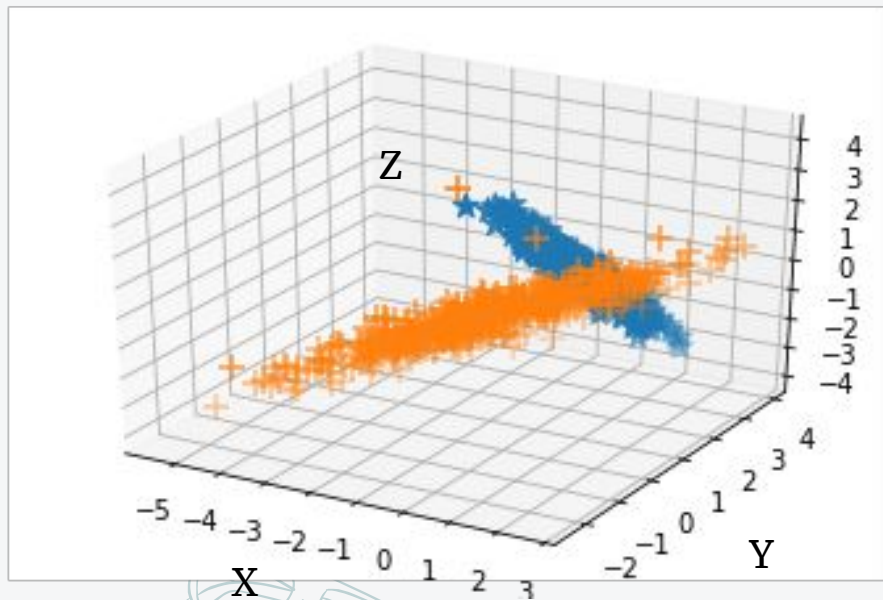
Number of features  $P = 3$



# DATA

Number of samples  $N = 1000$

Number of features  $P = 3$







# PRINCIPAL COMPONENT ANALYSIS

Principle components = axes of greatest variability

When we talk about  $X_1, X_2, \dots, X_p$  – these features are in cartesian coordinate system

Principle components are a *new* set of  $p$  axes in the direction of greatest variability, that are a linear combination of the original



## SELECTING PRINCIPAL COMPONENTS

In one dimension, we would projecting points to a line.

How do we do that?



## SELECTING PRINCIPAL COMPONENTS

In one dimension, we would project points to a line.

How do we do that? Regression (squared loss)

How do we do it in many dimensions?

# LINEAR ALGEBRA (REVIEW)

- Eigenvectors for a square  $m \times m$  matrix  $S$

$$Sv = \lambda v$$

Eigenvector  $v \in \mathbb{R}^m, v \neq 0$

Eigenvalue  $\lambda \in \mathbb{R}$

E.g

$$\begin{pmatrix} 6 & -2 \\ 4 & 0 \end{pmatrix} \underbrace{\begin{pmatrix} 1 \\ 2 \end{pmatrix}}_{\text{eigenvector}} = \begin{pmatrix} 2 \\ 4 \end{pmatrix} = \underbrace{2}_{\text{eigenvalue}} \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$



# LINEAR ALGEBRA (REVIEW)

How many eigenvalues are there at most?



## LINEAR ALGEBRA (REVIEW)

For *symmetric matrices*, eigenvectors for distinct eigenvalues are *orthogonal*

All eigenvalues for a *positive semidefinite matrix* are *non-negative*

Note:  $S$  is positive semidefinite if  $\forall w \in \mathbb{R}^m, w^T S w \geq 0$

# EIGEN DECOMPOSITION

Eigen decomposition transforms a matrix into its principal vectors of variation

$$S = \begin{array}{c} u_1, u_2, u_3, \dots, u_m \\ \begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array} \\ U \end{array} \begin{array}{c} \lambda_1 \\ \lambda_2 \\ \\ \lambda_m \end{array} U^{-1}$$

$$S = U \Lambda U^{-1}$$

U are eigenvectors of S, diagonal values of Lambda are eigenvalues.

# EIGEN DECOMPOSITION

Eigen decomposition transforms a matrix into its principal vectors of variation

$$S = \begin{matrix} u_1, u_2, u_3, \dots, u_m \\ \boxed{\begin{array}{|c|c|c|c|c|} \hline & & & & \\ \hline \end{array}} U \begin{array}{|c|c|} \hline \lambda_1 & 0 \\ \hline 0 & \lambda_m \\ \hline \end{array} \boxed{U^{-1}} \end{matrix}$$

$$S = U \Lambda U^{-1}$$

Decomposition is unique if eigenvalues are unique

Columns of U are **eigenvectors** of S

Diagonal values of Lambda are **eigenvalues** of S



# SINGULAR VALUE DECOMPOSITION (SVD)

For an **m x n** matrix **A** of **rank r** there exists a factorization (SVD) as follows:

$$A = U \Sigma V^T$$

$\begin{matrix} \nearrow & \nearrow & \nearrow \\ m \times m & m \times n & n \times n \end{matrix}$

The columns of  $U$  are orthogonal eigenvectors of  $\mathbf{A}\mathbf{A}^T$

The columns of  $V$  are orthogonal eigenvectors of  $\mathbf{A}^T\mathbf{A}$

Eigenvalues  $(\sigma_1, \sigma_2, \dots, \sigma_r)$  of  $\mathbf{A}\mathbf{A}^T$  are eigenvalues of  $\mathbf{A}^T\mathbf{A}$

$$\Sigma = \begin{bmatrix} \sigma_1 & & 0 \\ & \sigma_2 & \\ 0 & & \sigma_r \end{bmatrix}$$

Note: rank  $r$  means there are max of  $r$  linearly independent rows and columns in  $A$

# SINGULAR VALUE DECOMPOSITION (SVD)

For an **m x n** matrix **A** of **rank r** there exists a factorization (SVD) as follows:

$$A = U \Sigma V^T$$

$\begin{matrix} \nearrow & \nearrow & \nwarrow \\ m \times m & m \times n & n \times n \end{matrix}$

The columns of  $U$  are orthogonal eigenvectors of  $AA^T$

The columns of  $V$  are orthogonal eigenvectors of  $A^T A$

Eigenvalues  $(\sigma_1, \sigma_2, \dots, \sigma_r)$  of  $AA^T$  are eigenvalues of  $A^T A$

$$\Sigma = \begin{bmatrix} \sigma_1 & & 0 \\ & \sigma_2 & \\ 0 & & \sigma_r \end{bmatrix} \leftarrow \begin{matrix} \text{Singular} \\ \text{Values} \end{matrix}$$

Note: rank  $r$  means there are max of  $r$  linearly independent rows and columns in  $A$

# SINGULAR VALUE DECOMPOSITION (SVD)

For an **m x n** matrix **A** of **rank r** there exists a factorization (SVD) as follows:

$$A = U \Sigma V^T$$

$\begin{matrix} \nearrow & \nearrow & \nwarrow \\ m \times m & m \times n & n \times n \end{matrix}$

The columns of  $U$  are orthogonal eigenvectors of  $\mathbf{A}\mathbf{A}^T$

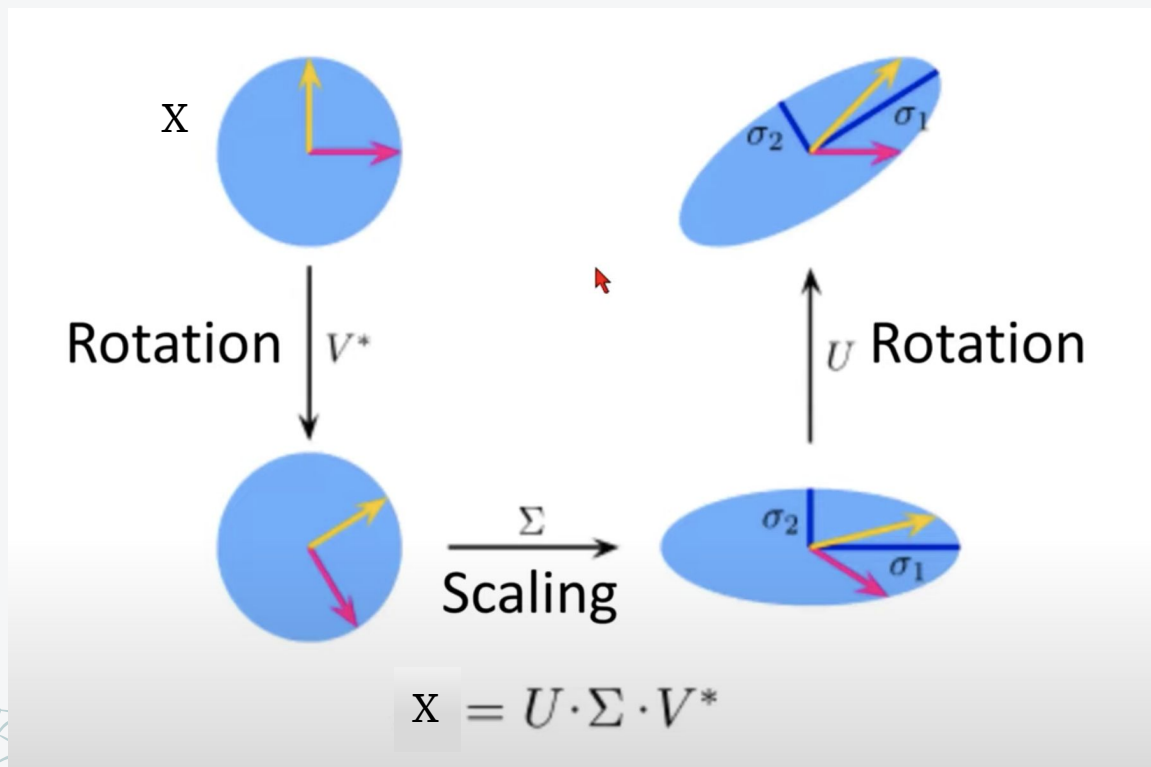
The columns of  $V$  are orthogonal eigenvectors of  $\mathbf{A}^T\mathbf{A}$

Eigenvalues  $(\sigma_1, \sigma_2, \dots, \sigma_r)$  of  $\mathbf{A}\mathbf{A}^T$  are eigenvalues of  $\mathbf{A}^T\mathbf{A}$

$$\Sigma = \begin{bmatrix} \sigma_1 & & 0 \\ & \sigma_2 & \\ 0 & & \sigma_r \end{bmatrix} \quad \leftarrow \quad \lambda_i = \sqrt{\sigma_i}$$

Note: rank  $r$  means there are max of  $r$  linearly independent rows and columns in  $A$

# GEOMETRIC INTERPRETATION OF SVD



## LOW RANK APPROXIMATION

Rank  $r$  of matrix  $A$  might not be small, but we can find  $A_k$  with rank  $k \ll r$  :

$$A_k = \min_{Z: \text{rank}(Z)=k} \|A - Z\|_F$$

$A$  and  $Z$  are both  $m \times n$  matrices

Frobenius  
(Euclidean) norm

$$\|X\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |X_{ij}|^2}$$

# LOW RANK APPROXIMATION

## Solution via SVD

$$A_k = U \text{diag}(\sigma_1, \dots, \sigma_k, \underbrace{0, \dots, 0}_{\text{Set smallest } r-k \text{ singular values to } 0}) V^T$$

Error:  $\|A - A_k\|_F = \sigma_{k+1}$



## BACK TO PCA

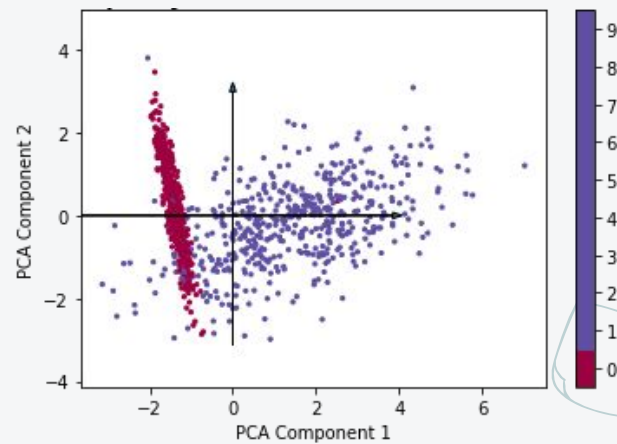
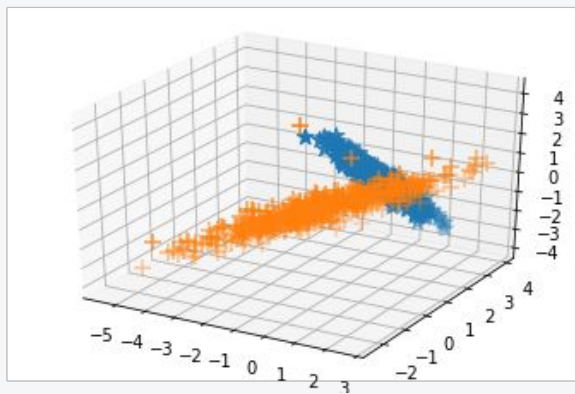
- Find eigenvectors and eigenvalues using SVD
- Examine how much variance is explained using  $k$  components ( $k$  is varied)
- For the purposes of explaining (reconstructing the data), pick the number of components based on variance explained
- For plotting purposes can plot the first 2 PCA or all pairs of interest

PCA

Variance explained:

PC1	PC2	PC3
[0.73	0.24	0.03]

DATA

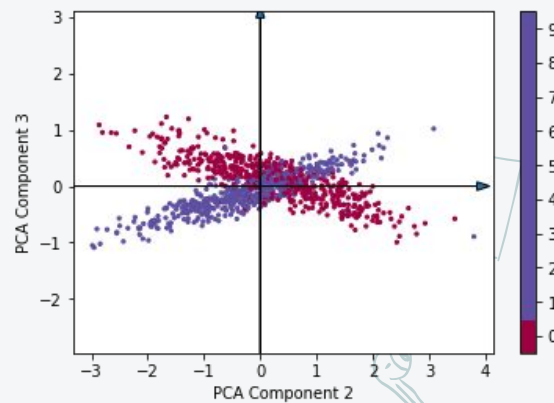
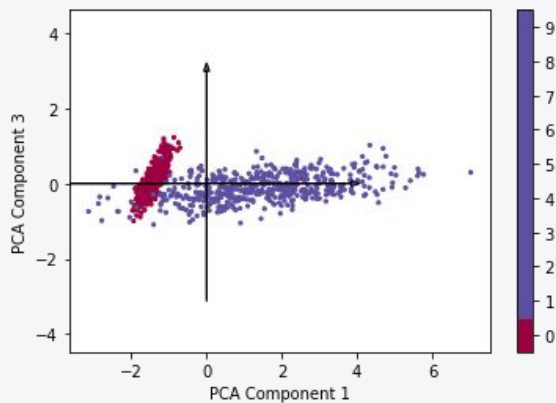
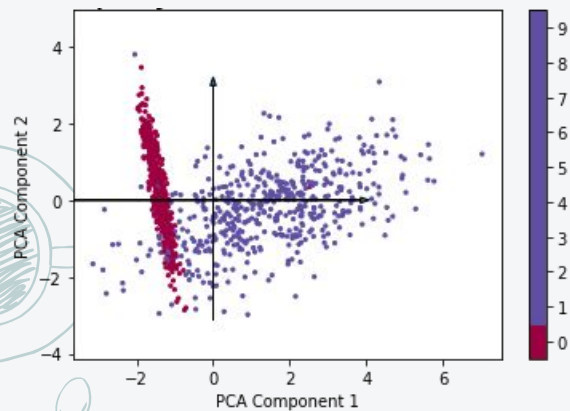




# PCA

Variance explained:

PC1	PC2	PC3
[0.73	0.24	0.03]



## ANOTHER DATASET: DIGITS

0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9
8	9	1	7	7	3	5	1	0	0	2	2	7	7	2	0	1	2	6	3
3	2	3	3	4	6	6	6	4	7	1	5	0	5	5	2	2	2	0	0
1	7	6	3	2	1	7	9	6	3	1	3	3	1	7	6	8	4	3	1
4	0	5	7	6	5	6	1	7	5	6	4	7	2	2	2	2	5	7	5
5	9	8	1	9	9	0	7	9	1	0	1	2	3	4	5	6	7	8	9
0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9
0	3	5	5	6	5	0	9	8	5	2	4	1	7	7	3	5	1	0	0
2	2	7	8	1	0	1	2	6	3	3	7	3	3	4	6	6	6	4	9
1	5	0	9	5	2	3	1	0	0	1	7	6	3	2	1	7	3	1	3
9	1	7	6	8	4	3	1	4	0	5	7	6	3	6	1	7	5	4	4
7	1	2	1	1	5	5	4	8	8	4	9	0	8	9	8	0	1	2	3
4	5	6	7	8	7	0	1	2	3	4	5	6	7	8	9	0	1	2	3
4	5	6	7	8	9	0	3	5	5	6	5	0	9	8	9	8	4	1	7
7	3	5	1	0	0	2	2	7	8	2	0	1	2	6	3	3	7	3	3
4	6	6	6	4	7	1	5	0	9	5	2	8	2	0	0	1	7	6	3
2	1	2	4	6	3	1	3	9	1	7	6	8	4	3	1	9	0	5	3
6	7	6	1	7	5	4	4	7	2	8	2	2	5	7	9	5	4	8	8
4	9	0	8	9	3	0	1	1	3	4	5	6	7	8	9	0	1	2	3

Optical recognition of handwritten digits dataset

Number of Instances: 5620

Number of Attributes: 64

Attribute Information: 8x8 image of integer pixels in the range 0..16

Creator: E. Alpaydin

Date: July, 1998

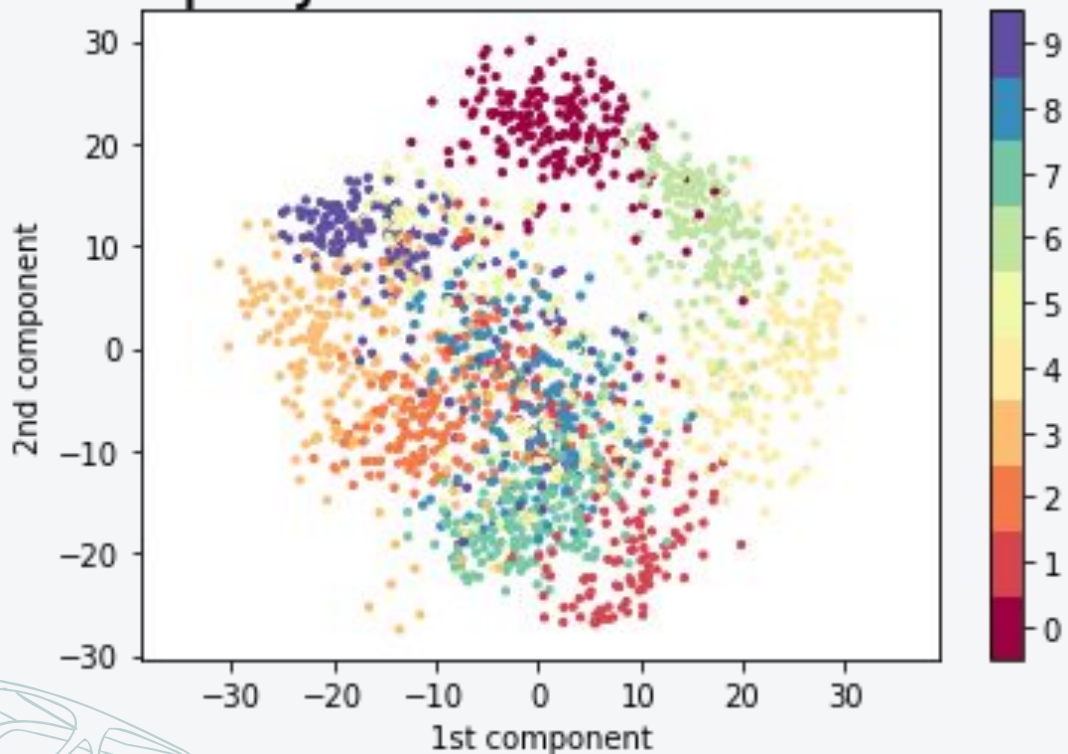
0	1	4	2	7	4	5	4	7	8	9	0	4	1	5	6	8	7	5		
0	1	4	2	7	4	5	4	7	8	9	0	4	1	5	6	8	7	5		
0	2	9	1	7	7	3	5	1	0	0	2	5	7	2	0	1	2	6	3	1
3	7	3	3	4	6	6	6	6	4	7	1	5	0	5	5	2	1	2	0	0
1	7	4	2	7	4	5	4	7	8	9	0	4	1	5	6	8	7	5	4	0
4	0	5	7	4	5	4	7	8	9	0	4	1	5	6	8	7	5	4	0	0
0	3	5	6	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1
1	2	3	4	5	6	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0
4	5	0	9	1	2	3	4	5	6	7	8	9	0	4	1	2	3	4	5	6
9	4	7	8	9	2	0	1	3	1	0	1	1	2	3	4	5	6	7	8	9
7	1	8	1	1	5	5	4	8	8	4	9	0	8	9	8	0	4	2	3	0
9	5	6	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1	2
6	5	4	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1	2
7	3	5	4	0	0	2	2	7	9	2	0	4	2	6	3	3	7	3	3	3
4	6	6	6	6	6	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0
2	4	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1	2	3
6	7	8	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1	2	3	4
9	9	0	4	1	2	3	4	5	6	7	8	9	0	4	1	2	3	4	5	6



# PCA: DIGITS

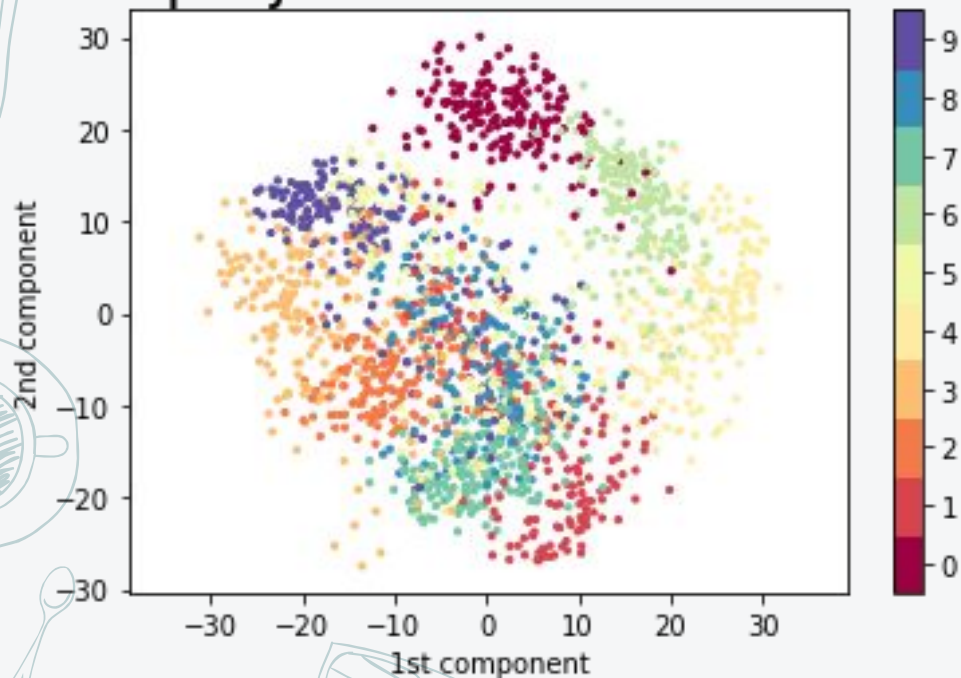


## PCA projection of the dataset

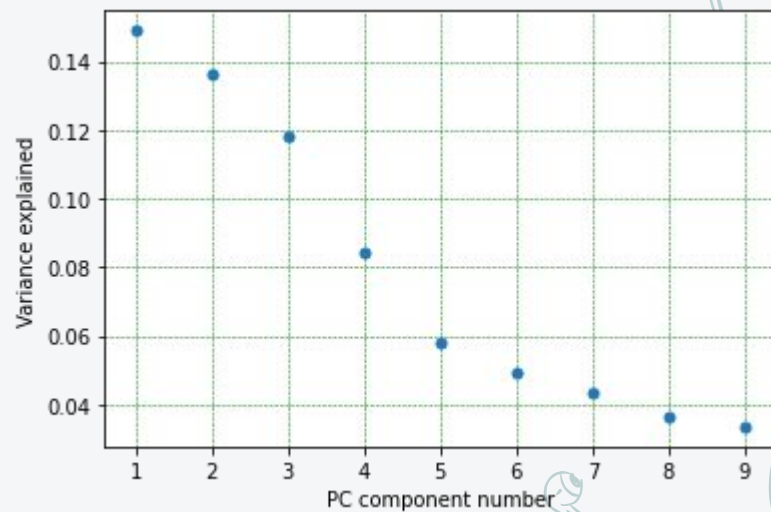


# PCA: DIGITS

## PCA projection of the dataset



## Variance explained





# PCA

## Advantages:

- Relatively Fast
- Gives some sense of intrinsic dimensionality
- Can compute variation explained
- Axes are interpretable

## Disadvantages:

- Linear reduction in dimensionality
- Orthogonal components (hard to model non-linear data)



# NON-LINEAR LOWER DIMENSIONAL EMBEDDING

Neighborhood based

- T-SNE
- UMAP
- ISOMAP
- LLE

Autoencoder-based

And many more...



## T-SNE : T-DISTRIBUTED STOCHASTIC NETWORK EMBEDDING

- *Distance preservation*: when 2 points are close in high dimensional space, they should be close in 2 dimensional space
- *Neighbor preservation*: neighborhoods of a given point in high-d should be neighborhoods of a given point in low-d

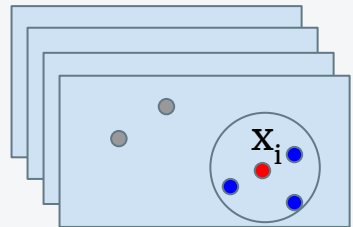
Note1: effective number of neighbors is a parameter called *perplexity*

Note2: t-SNE focuses on the order of things, not necessarily on the distance itself

# T-SNE

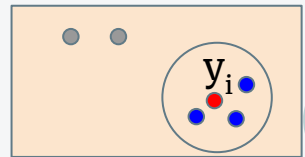
Similarity of data points in High Dimensional space

$$p_{ij} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma^2)}{\sum_{k \neq l} \exp(-\|x_k - x_l\|^2 / 2\sigma^2)}$$



Similarity of points in Low Dimensional space

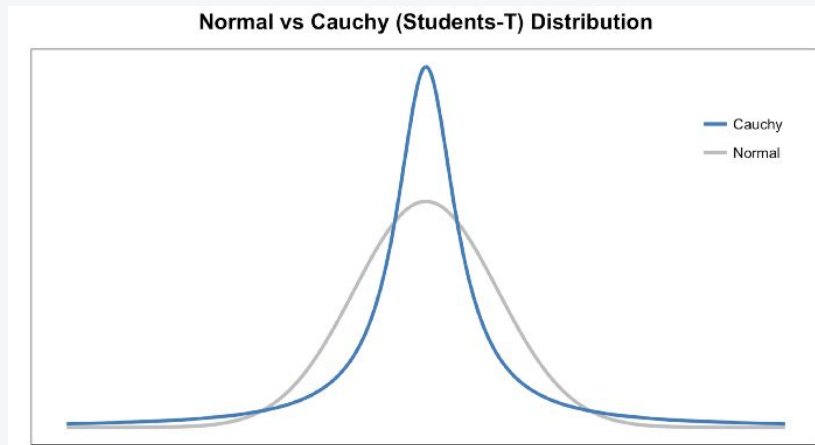
$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$





# T-SNE

Reason for t-distribution as opposed to Gaussian as in high  $d$  is crowding – need fatter tails to accommodate distant points



# T-SNE COST FUNCTION

Cost function: Kullback Leibler divergence (relative entropy)  
- measures distance between two distributions

$$C = KL(P||Q) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

(Kullback and Leibler, 1951)

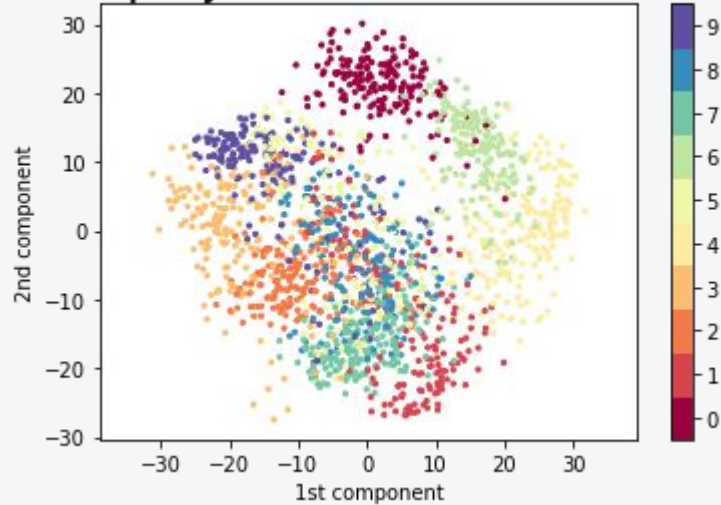
Properties?



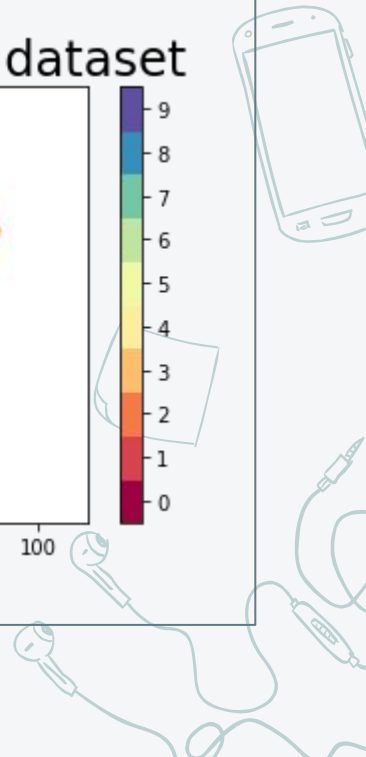
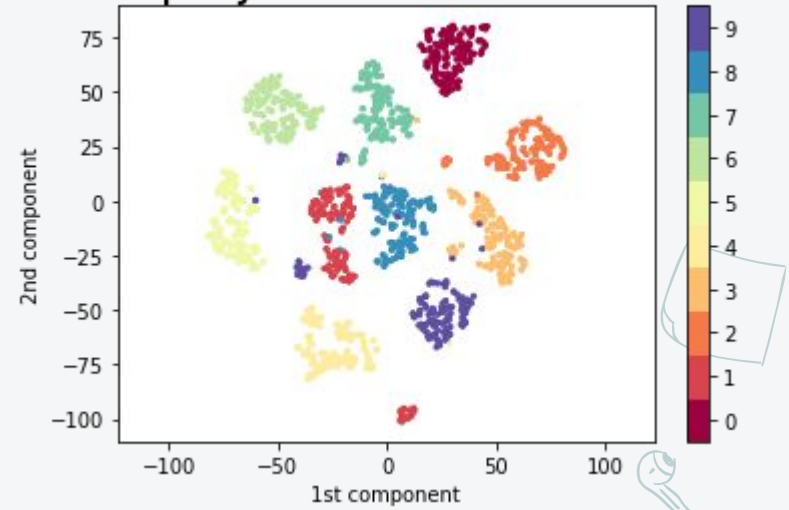
# T-SNE ON DIGITS



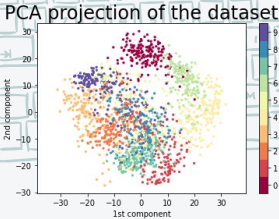
PCA projection of the dataset



TSNE projection of the dataset



PCA projection of the dataset



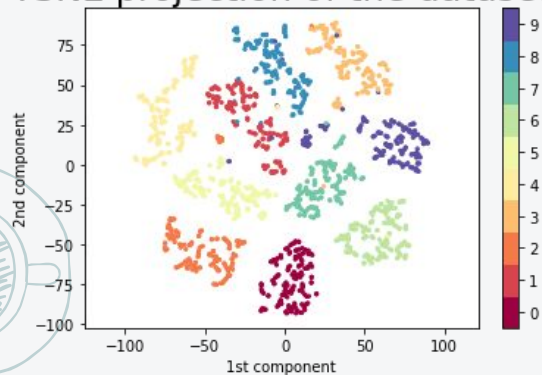
# T-SNE SENSITIVITY TO PARAMETERS: PERPLEXITY

Perplexity = 5

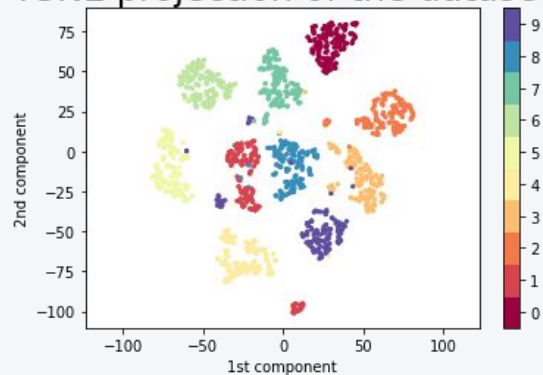
Perplexity = 10

Perplexity = 50

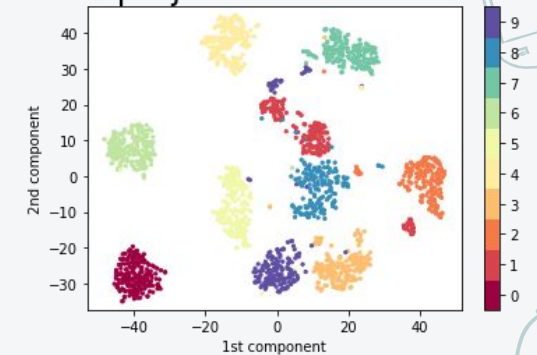
TSNE projection of the dataset



TSNE projection of the dataset

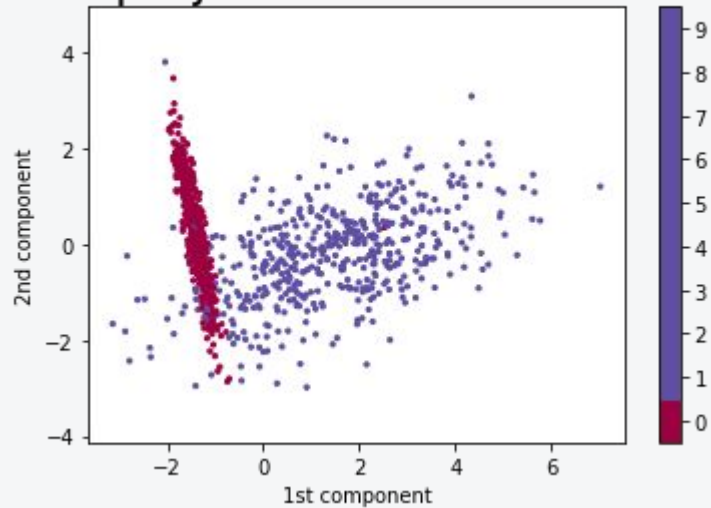


TSNE projection of the dataset

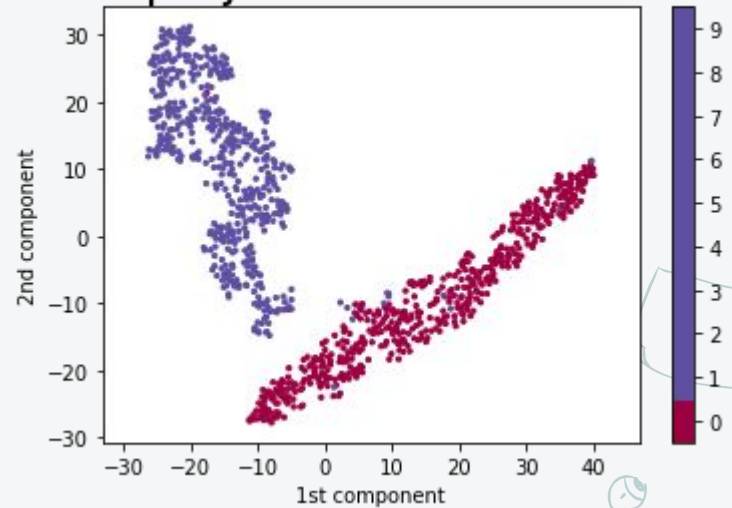


# T-SNE ON THE ORIGINAL DATASET

PCA projection of the dataset



TSNE projection of the dataset



# T-SNE

## Advantages

- Retains local structure
- Excels at visualizing complex high-d data in 2d

## Disadvantages

- Slower
- Doesn't really capture global structure
- Sensitive to params
- Clustering and distance depend on perplexity (shouldn't be interpreted off the plot)
- Stochasticity leads to different results
- Not guaranteed to converge to the global optimum of the cost function



# UMAP: UNIFORM MANIFOLD APPROXIMATION AND PROJECTION

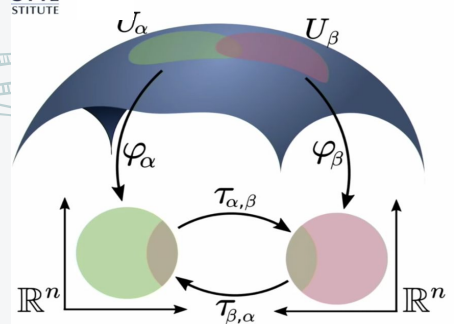
UMAP – preserves both local and global structure by

1. Learning the manifold of the data
2. Embedding that manifold into a low dimensional space
  - Builds on neighborhood based approaches but adds mathematical foundations

Key intuition:

While the assumptions may not be true in the original space, they can be true on the manifold

# UMAP



## Assumptions:

- The points are covering manifold uniformly
- There is local connectedness between points (confident in the distance to the 1st neighbor)

Using topological theory, construct the neighborhood graph on the low dimensional manifold.

Optimize cross entropy:

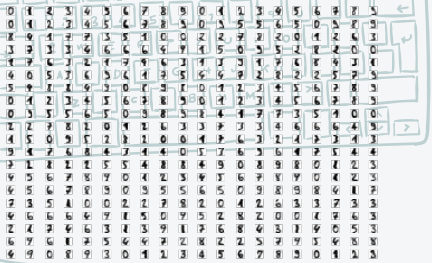
Get the clumps right

Similar to t-SNE

$$\sum_{a \in A} \mu(a) \log \left( \frac{\mu(a)}{\nu(a)} \right) + (1 - \mu(a)) \log \left( \frac{1 - \mu(a)}{1 - \nu(a)} \right)$$

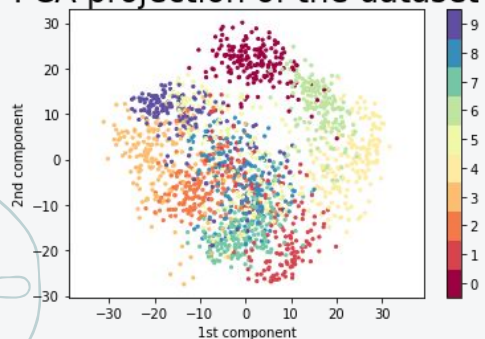
Get the gaps right



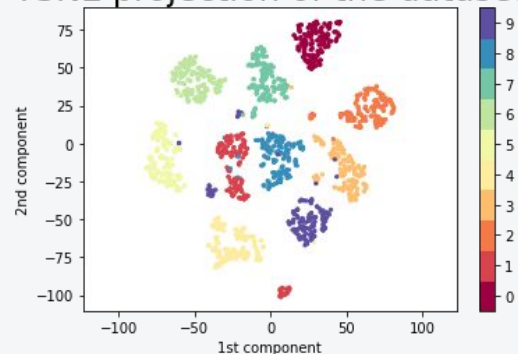


# UMAP ON DIGITS

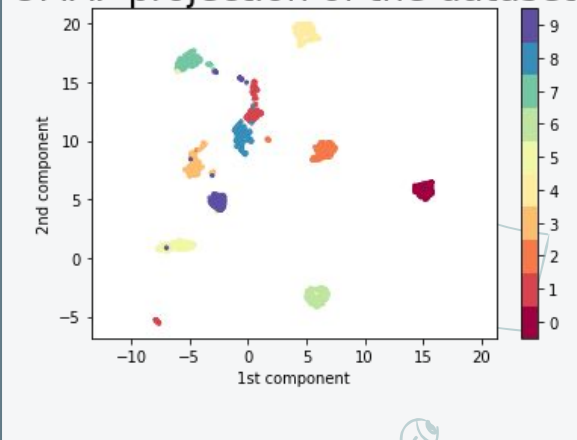
PCA projection of the dataset



t-SNE projection of the dataset



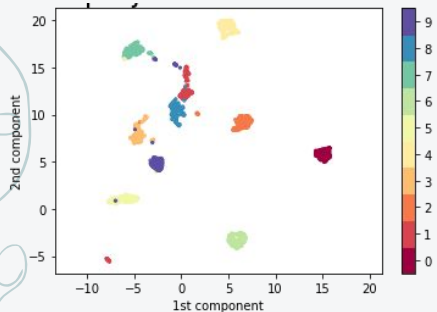
UMAP projection of the dataset



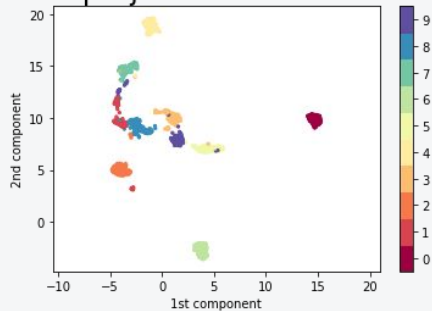
# UMAP HYPERPARAMETERS

- Number of neighbors - local vs global structure
- Minimum distance - how tightly points can be packed together
- Number of components (we set 2 for all)
- Distance metric

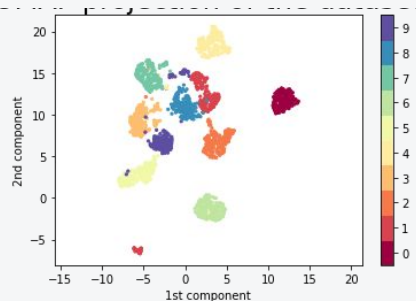
neighbors = 15  
metric = euclidian  
min\_dist = 0.1



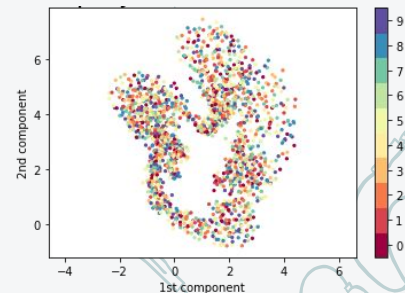
neighbors = 50  
metric = euclidian  
min\_dist = 0.1



neighbors = 5  
metric = euclidian  
min\_dist = 0.5



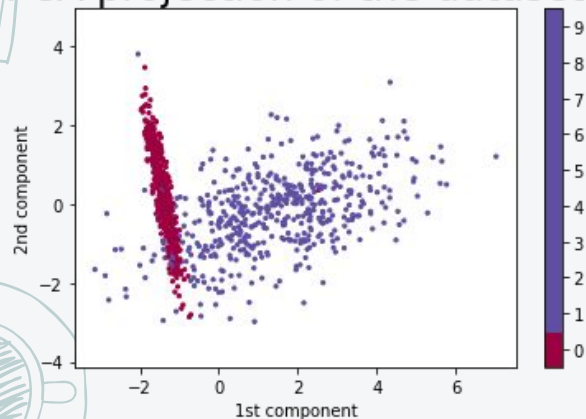
neighbors = 5  
metric = mahalanobis  
min\_dist = 0.1



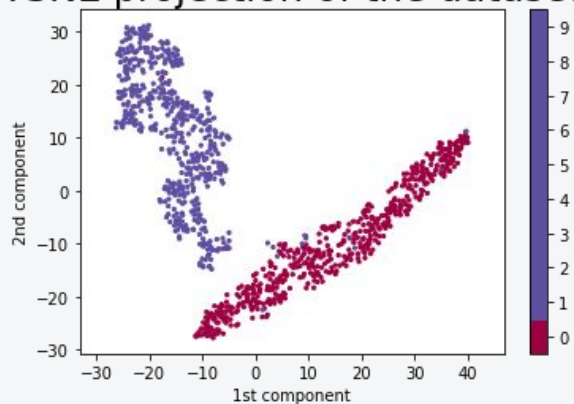
<https://umap-learn.readthedocs.io/en/latest/parameters.html> - on all parameters and their effects

# UMAP ON THE ORIGINAL DATA

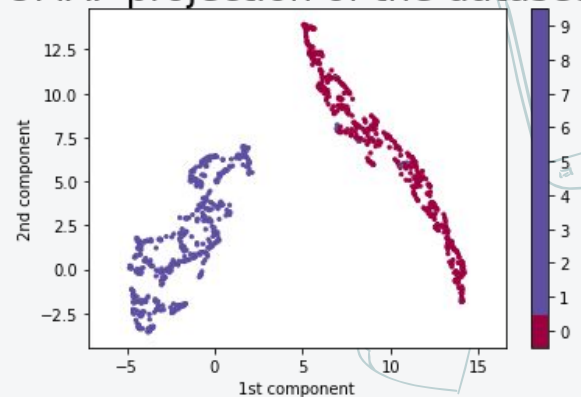
PCA projection of the dataset



TSNE projection of the dataset



UMAP projection of the dataset





# UMAP

## Advantages

- Preserves global as well as local structure
- Has mathematical guarantees
- More robust to hyperparameter changes (e.g. number of neighbors)

## Disadvantages

- Depends on hyperparameters
- Slow-ish

# READING

## PCA

- In depth PCA with examples from the Python handbook:  
<https://jakevdp.github.io/PythonDataScienceHandbook/05.09-principal-component-analysis.html>

## t-SNE

- Playing with t-sne: <https://distill.pub/2016/misread-tsne/>
- Intro to t-sne with python:  
<https://towardsdatascience.com/an-introduction-to-t-sne-with-python-example-5a3a293108d1>

## UMAP

- Talk (26min): <https://www.youtube.com/watch?v=nq6iPZVUxZU>
- Examples with code in python:  
[https://umap-learn.readthedocs.io/en/latest/basic\\_usage.html](https://umap-learn.readthedocs.io/en/latest/basic_usage.html)

## t-SNE vs UMAP:

<https://towardsdatascience.com/tsne-vs-umap-global-structure-4d8045acba17>